

Linear algebra in R

Søren Højsgaard

February 15, 2005

Contents

1	Introduction	1
2	Vectors	1
2.1	Vectors	1
2.2	Transpose of vectors	2
2.3	Multiplying a vector by a number	3
2.4	Sum of vectors	3
2.5	(Inner) product of vectors	4
2.6	The length (norm) of a vector	5
2.7	The 0-vector and 1-vector	5
2.8	Orthogonal (perpendicular) vectors	5
3	Matrices	6
3.1	Matrices	6
3.2	Multiplying a matrix with a number	6
3.3	Transpose of matrices	7
3.4	Sum of matrices	7
3.5	Multiplication of a matrix and a vector	7
3.6	Multiplication of matrices	8
3.7	Vectors as matrices	9
3.8	Some special matrices	9
3.9	Inverse of matrices	10
3.10	Solving systems of linear equations	11
3.11	Trace	12
3.12	Determinant	12
3.13	Some additional rules for matrix operations	12
3.14	Details on inverse matrices*	12
3.14.1	Inverse of a 2×2 matrix*	12
3.14.2	Inverse of diagonal matrices*	13
3.14.3	Generalized inverse*	13
3.14.4	Inverting an $n \times n$ matrix*	13
4	Least squares	15
5	A neat little exercise – from a bird’s perspective	16

1 Introduction

This note has two goal: 1) Introducing linear algebra (vectors and matrices) and 2) showing how to work with these concepts in R.

2 Vectors

2.1 Vectors

A column vector is a list of numbers stacked on top of each other, e.g.

$$a = \begin{bmatrix} 2 \\ 1 \\ 3 \end{bmatrix}$$

A row vector is a list of numbers written one after the other, e.g.

$$b = (2, 1, 3)$$

In both cases, the list is ordered, i.e.

$$(2, 1, 3) \neq (1, 2, 3).$$

We make the following convention:

- In what follows all vectors are column vectors unless otherwise stated.
- However, writing column vectors takes up more space than row vectors. Therefore we shall frequently write vectors as row vectors, but with the understanding that it really is a column vector.

A general n -vector has the form

$$a = \begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_n \end{bmatrix}$$

where the a_i s are numbers, and this vector shall be written $a = (a_1, \dots, a_n)$.

A graphical representation of 2-vectors is shown Figure 1. Note that row and




Figure 1: Two 2-vectors

column vectors are drawn the same way.

```
> a <- c(1, 3, 2)
> a
```

```
[1] 1 3 2
```

The vector `a` is in R printed “in row format” but can really be regarded as a column vector, cfr. the convention above.

2.2 Transpose of vectors

Transposing a vector means turning a column (row) vector into a row (column) vector. The transpose is denoted by “ \top ”.

Example 1

$$\begin{bmatrix} 1 \\ 3 \\ 2 \end{bmatrix}^\top = [1, 3, 2] \quad \text{og} \quad [1, 3, 2]^\top = \begin{bmatrix} 1 \\ 3 \\ 2 \end{bmatrix}$$

□

Hence transposing twice takes us back to where we started:

$$a = (a^\top)^\top$$

```
> t(a)
```

```
      [,1] [,2] [,3]
[1,]     1     3     2
```

2.3 Multiplying a vector by a number

If a is a vector and α is a number then αa is the vector

$$\alpha a = \begin{bmatrix} \alpha a_1 \\ \alpha a_2 \\ \vdots \\ \alpha a_n \end{bmatrix}$$

See Figure 2.

Example 2

$$7 \begin{bmatrix} 1 \\ 3 \\ 2 \end{bmatrix} = \begin{bmatrix} 7 \\ 21 \\ 14 \end{bmatrix}$$

□

```
> 7 * a
```

```
[1]  7 21 14
```




Figure 2: Multiplication of a vector by a number

2.4 Sum of vectors

Let a and b be n -vectors. The sum $a + b$ is the n -vector

$$a + b = \begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_n \end{bmatrix} + \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{bmatrix} = \begin{bmatrix} a_1 + b_1 \\ a_2 + b_2 \\ \vdots \\ a_n + b_n \end{bmatrix} = b + a$$

See Figure 3 and 4. Only vectors of the same dimension can be added.

Example 3

$$\begin{bmatrix} 1 \\ 3 \\ 2 \end{bmatrix} + \begin{bmatrix} 2 \\ 8 \\ 9 \end{bmatrix} = \begin{bmatrix} 1+2 \\ 3+8 \\ 2+9 \end{bmatrix} = \begin{bmatrix} 3 \\ 11 \\ 11 \end{bmatrix}$$

□




Figure 3: Addition of vectors




Figure 4: Addition of vectors and multiplication by a number

```
> a <- c(1, 3, 2)
> b <- c(2, 8, 9)
> a + b
```

```
[1] 3 11 11
```

2.5 (Inner) product of vectors

Let $a = (a_1, \dots, a_n)$ and $b = (b_1, \dots, b_n)$. The (inner) product of a and b is

$$a \cdot b = a_1 b_1 + \dots + a_n b_n$$

Note, that the product is a number – not a vector.

```
> sum(a * b)
```

```
[1] 44
```

2.6 The length (norm) of a vector

The length (or norm) of a vector a is

$$||a|| = \sqrt{a \cdot a} = \sqrt{\sum_{i=1}^n a_i^2}$$

```
> sqrt(sum(a * a))
```

```
[1] 3.741657
```

2.7 The 0–vector and 1–vector

The 0–vector (1–vector) is a vector with 0 (1) on all entries. The 0–vector (1–vector) is frequently written simply as 0 (1) or as 0_n (1_n) to emphasize that its length n .

```
> rep(0, 5)
```

```
[1] 0 0 0 0 0
```

```
> rep(1, 5)
```

```
[1] 1 1 1 1 1
```

2.8 Orthogonal (perpendicular) vectors

Two vectors v_1 and v_2 are orthogonal if their inner product is zero, written

$$v_1 \perp v_2 \Leftrightarrow v_1 \cdot v_2 = 0$$

```
> v1 <- c(1, 1)
> v2 <- c(-1, 1)
> sum(v1 * v2)
```

```
[1] 0
```

3 Matrices

3.1 Matrices

An $r \times c$ matrix A (reads “an r times c matrix”) is a table with r rows og c columns

$$A = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1c} \\ a_{21} & a_{22} & \cdots & a_{2c} \\ \vdots & \vdots & \ddots & \vdots \\ a_{r1} & a_{r2} & \cdots & a_{rc} \end{bmatrix}$$

Note that one can regard A as consisting of c columns vectors put after each other:

$$A = [a_1 : a_2 : \cdots : a_c]$$

```
> A <- matrix(c(1, 3, 2, 2, 8, 9), ncol = 3)
> A
```

```
      [,1] [,2] [,3]
[1,]     1     2     8
[2,]     3     2     9
```

Note that the numbers 1,3,2,2,8,9 are read into the matrix column-by-column. To get the numbers read in row-by-row do

```
> A2 <- matrix(c(1, 3, 2, 2, 8, 9), ncol = 3, byrow = T)
> A2
```

```
      [,1] [,2] [,3]
[1,]    1    3    2
[2,]    2    8    9
```

3.2 Multiplying a matrix with a number

For a number α and a matrix A , the product αA is the matrix obtained by multiplying each element in A by α .

Example 4

$$7 \begin{bmatrix} 1 & 2 \\ 3 & 8 \\ 2 & 9 \end{bmatrix} = \begin{bmatrix} 7 & 14 \\ 21 & 56 \\ 14 & 63 \end{bmatrix}$$

□

```
> 7 * A
```

```
      [,1] [,2] [,3]
[1,]    7   14   56
[2,]   21   14   63
```

3.3 Transpose of matrices

A matrix is transposed by interchanging rows and columns and is denoted by “ \top ”.

Example 5

$$\begin{bmatrix} 1 & 2 \\ 3 & 8 \\ 2 & 9 \end{bmatrix}^{\top} = \begin{bmatrix} 1 & 3 & 2 \\ 2 & 8 & 9 \end{bmatrix}$$

□

Note that if A is an $r \times c$ matrix then A^{\top} is a $c \times r$ matrix.

```
> t(A)
```

```
      [,1] [,2]
[1,]    1    3
[2,]    2    2
[3,]    8    9
```

3.4 Sum of matrices

Let A and B be $r \times c$ matrices. The sum $A + B$ is the $r \times c$ matrix obtained by adding A and B elementwise.

Only matrices with the same dimensions can be added.

Example 6

$$\begin{bmatrix} 1 & 2 \\ 3 & 8 \\ 2 & 9 \end{bmatrix} + \begin{bmatrix} 5 & 4 \\ 8 & 2 \\ 3 & 7 \end{bmatrix} = \begin{bmatrix} 6 & 6 \\ 11 & 10 \\ 5 & 16 \end{bmatrix}$$

□

```
> B <- matrix(c(5, 8, 3, 4, 2, 7), ncol = 3, byrow = T)
> A + B
```

```
      [,1] [,2] [,3]
[1,]     6    10    11
[2,]     7     4    16
```

3.5 Multiplication of a matrix and a vector

Let A be an $r \times c$ matrix and let b be a c -dimensional column vector. The product Ab is the $r \times 1$ matrix

$$Ab = \begin{bmatrix} a_{11} & a_{12} & \dots & a_{1c} \\ a_{21} & a_{22} & \dots & a_{2c} \\ \vdots & \vdots & \ddots & \vdots \\ a_{r1} & a_{r2} & \dots & a_{rc} \end{bmatrix} \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_c \end{bmatrix} = \begin{bmatrix} a_{11}b_1 + a_{12}b_2 + \dots + a_{1c}b_c \\ a_{21}b_1 + a_{22}b_2 + \dots + a_{2c}b_c \\ \vdots \\ a_{r1}b_1 + a_{r2}b_2 + \dots + a_{rc}b_c \end{bmatrix}$$

Example 7

$$\begin{bmatrix} 1 & 2 \\ 3 & 8 \\ 2 & 9 \end{bmatrix} \begin{bmatrix} 5 \\ 8 \end{bmatrix} = \begin{bmatrix} 1 \cdot 5 + 2 \cdot 8 \\ 3 \cdot 5 + 8 \cdot 8 \\ 2 \cdot 5 + 9 \cdot 8 \end{bmatrix} = \begin{bmatrix} 21 \\ 79 \\ 82 \end{bmatrix}$$

□

```
> A %*% a
```

```
      [,1]
[1,]    23
[2,]    27
```

Note the difference to

```
> A * a
```

```
      [,1] [,2] [,3]
[1,]     1     4    24
[2,]     9     2    18
```

Figure out yourself what goes on!

3.6 Multiplication of matrices

Let A be an $r \times c$ matrix and B a $c \times t$ matrix, i.e. $B = [b_1 : b_2 : \dots : b_t]$. The product AB is the $r \times t$ matrix given by:

$$AB = A[b_1 : b_2 : \dots : b_t] = [Ab_1 : Ab_2 : \dots : Ab_t]$$

Example 8

$$\begin{aligned} \begin{bmatrix} 1 & 2 \\ 3 & 8 \\ 2 & 9 \end{bmatrix} \begin{bmatrix} 5 & 4 \\ 8 & 2 \end{bmatrix} &= \left[\begin{bmatrix} 1 & 2 \\ 3 & 8 \\ 2 & 9 \end{bmatrix} \begin{bmatrix} 5 \\ 8 \end{bmatrix} : \begin{bmatrix} 1 & 2 \\ 3 & 8 \\ 2 & 9 \end{bmatrix} \begin{bmatrix} 4 \\ 2 \end{bmatrix} \right] \\ &= \begin{bmatrix} 1 \cdot 5 + 2 \cdot 8 & 1 \cdot 4 + 2 \cdot 2 \\ 3 \cdot 5 + 8 \cdot 8 & 3 \cdot 4 + 8 \cdot 2 \\ 2 \cdot 5 + 9 \cdot 8 & 2 \cdot 4 + 9 \cdot 2 \end{bmatrix} = \begin{bmatrix} 21 & 8 \\ 79 & 28 \\ 82 & 26 \end{bmatrix} \end{aligned}$$

□

Note that the product AB can only be formed if the number of rows in B and the number of columns in A are the same. In that case, A and B are said to be conforme.

In general AB and BA are not identical.

A [MNEMONIC FOR MATRIX MULTIPLICATION](#) is :

$$\begin{bmatrix} 1 & 2 \\ 3 & 8 \\ 2 & 9 \end{bmatrix} \begin{bmatrix} 5 & 4 \\ 8 & 2 \end{bmatrix} = \begin{array}{cc|cc} & & 5 & 4 \\ & & 8 & 2 \\ \hline 1 & 2 & 1 \cdot 5 + 2 \cdot 8 & 1 \cdot 4 + 2 \cdot 2 \\ 3 & 8 & 3 \cdot 5 + 8 \cdot 8 & 3 \cdot 4 + 8 \cdot 2 \\ 2 & 9 & 2 \cdot 5 + 9 \cdot 8 & 2 \cdot 4 + 9 \cdot 2 \end{array} = \begin{bmatrix} 21 & 8 \\ 79 & 28 \\ 82 & 26 \end{bmatrix}$$

```
> A <- matrix(c(1, 3, 2, 2, 8, 9), ncol = 2)
> B <- matrix(c(5, 8, 4, 2), ncol = 2)
> A %*% B
```

```
      [,1] [,2]
[1,]   21    8
[2,]   79   28
[3,]   82   26
```

3.7 Vectors as matrices

One can regard a column vector of length r as an $r \times 1$ matrix and a row vector of length c as a $1 \times c$ matrix.

3.8 Some special matrices

- An $n \times n$ matrix is a [SQUARE MATRIX](#)
- A matrix A is [SYMMETRIC](#) if $A = A^\top$.
- A matrix with 0 on all entries is the [0-MATRIX](#) and is often written simply as 0.

- A matrix consisting of 1s in all entries is often written J .
- A square matrix with 0 on all off-diagonal entries and elements d_1, d_2, \dots, d_n on the diagonal is a **DIAGONAL MATRIX** and is often written $\text{diag}\{d_1, d_2, \dots, d_n\}$
- A diagonal matrix with 1s on the diagonal is called the **IDENTITY MATRIX** and is denoted I . The identity matrix satisfies that $IA = AI = A$.

- 0-matrix and 1-matrix

```
> matrix(0, nrow = 2, ncol = 3)
```

```
      [,1] [,2] [,3]
[1,]    0    0    0
[2,]    0    0    0
```

```
> matrix(1, nrow = 2, ncol = 3)
```

```
      [,1] [,2] [,3]
[1,]    1    1    1
[2,]    1    1    1
```

- Diagonal matrix and identity matrix

```
> diag(c(1, 2, 3))
```

```
      [,1] [,2] [,3]
[1,]    1    0    0
[2,]    0    2    0
[3,]    0    0    3
```

```
> diag(1, 3)
```

```
      [,1] [,2] [,3]
[1,]    1    0    0
[2,]    0    1    0
[3,]    0    0    1
```

Note what happens when `diag` is applied to a matrix:

```
> diag(diag(c(1, 2, 3)))
```

```
[1] 1 2 3
```

```
> diag(A)
```

```
[1] 1 8
```

3.9 Inverse of matrices

In general, the inverse of an $n \times n$ matrix A is the matrix B (which is also $n \times n$) which when multiplied with A gives the identity matrix I . That is,

$$AB = BA = I.$$

One says that B is A 's inverse and writes $B = A^{-1}$. Likewise, A is B 's inverse.

Example 9 Let

$$A = \begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix} \quad B = \begin{bmatrix} -2 & 1.5 \\ 1 & -0.5 \end{bmatrix}$$

Now $AB = BA = I$ so $B = A^{-1}$. □

Example 10 If A is a 1×1 matrix, i.e. a number, for example $A = 4$, then $A^{-1} = 1/4$. □

Some facts about inverse matrices are:

- Only square matrices can have an inverse, but not all square matrices have an inverse.
- When the inverse exists, it is unique.
- Finding the inverse of a large matrix A is numerically complicated (but computers do it for us).

In Section ?? the issue of matrix inversion is discussed in more detail.

Finding the inverse of a matrix in R is done using the `solve()` function:

```
> A <- matrix(c(1, 3, 2, 4), ncol = 2, byrow = T)
> A
```

```
      [,1] [,2]
[1,]     1     3
[2,]     2     4
```

```
> B <- solve(A)
> B
```

```
      [,1] [,2]
[1,]    -2  1.5
[2,]     1 -0.5
```

```
> A %*% B
```

```
      [,1] [,2]
[1,]     1     0
[2,]     0     1
```

3.10 Solving systems of linear equations

Example 11 Matrices are closely related to systems of linear equations. Consider the two equations

$$\begin{aligned}x_1 + 3x_2 &= 7 \\ 2x_1 + 4x_2 &= 10\end{aligned}$$

The system can be written in matrix form

$$\begin{bmatrix} 1 & 3 \\ 2 & 4 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \end{bmatrix} = \begin{bmatrix} 7 \\ 10 \end{bmatrix} \quad \text{i.e. } Ax = b$$

Since $A^{-1}A = I$ and since $Ix = x$ we have

$$x = A^{-1}b = \begin{bmatrix} -2 & 1.5 \\ 1 & -0.5 \end{bmatrix} \begin{bmatrix} 7 \\ 10 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \end{bmatrix}$$

A geometrical approach to solving these equations is as follows: Isolate x_2 in the equations:

$$x_2 = \frac{7}{3} - \frac{1}{3}x_1 \quad x_2 = \frac{1}{4}5 - \frac{2}{4}x_1$$

These two lines are shown in Figure 5 from which it can be seen that the solution is $x_1 = 1, x_2 = 2$.




Figure 5: Solving two equations with two unknowns.

From the Figure it follows that there are 3 possible cases of solutions to the system

1. Exactly one solution – when the lines intersect in one point
2. No solutions – when the lines are parallel but not identical
3. Infinitely many solutions – when the lines coincide.

□

```
> A <- matrix(c(1, 2, 3, 4), ncol = 2)
> b <- c(7, 10)
> x <- solve(A) %*% b
> x
```

```
      [,1]
[1,]    1
[2,]    2
```

3.11 Trace

Missing

3.12 Determinant

Missing

3.13 Some additional rules for matrix operations

For matrices A, B and C whose dimension match appropriately: the following rules apply

$$(A + B)^{\top} = A^{\top} + B^{\top}$$

$$(AB)^{\top} = B^{\top} A^{\top}$$

$$A(B + C) = AB + AC$$

$$AB = AC \not\Rightarrow B = C$$

In general $AB \neq BA$

$$AI = IA = A$$

If α is a number then $\alpha AB = A(\alpha B)$

3.14 Details on inverse matrices*

3.14.1 Inverse of a 2×2 matrix*

It is easy find the inverse for a 2×2 matrix. When

$$A = \begin{bmatrix} a & b \\ c & d \end{bmatrix}$$

then the inverse is

$$A^{-1} = \frac{1}{ad - bc} \begin{bmatrix} d & -b \\ -c & a \end{bmatrix}$$

under the assumption that $ad - bc \neq 0$. The number $ad - bc$ is called the determinant of A , sometimes written $|A|$. If $|A| = 0$, then A has no inverse.

3.14.2 Inverse of diagonal matrices*

Finding the inverse of a diagonal matrix is easy: Let

$$A = \text{diag}(a_1, a_2, \dots, a_n)$$

where all $a_i \neq 0$. Then the inverse is

$$A^{-1} = \text{diag}\left(\frac{1}{a_1}, \frac{1}{a_2}, \dots, \frac{1}{a_n}\right)$$

If one $a_i = 0$ then A^{-1} does not exist.

3.14.3 Generalized inverse*

Not all square matrices have an inverse. However all square matrices have an infinite number of generalized inverses. A generalized inverse of a square matrix A is a matrix A^{-} satisfying that

$$AA^{-}A = A.$$

For many practical problems it suffice to find a generalized inverse.

3.14.4 Inverting an $n \times n$ matrix*

In the following we will illustrate one frequently applied method for matrix inversion. The method is called Gauss-Seidel's method and many computer programs, including `solve()` use variants of the method for finding the inverse of an $n \times n$ matrix.

Consider the matrix A :

```
> A <- matrix(c(2, 2, 3, 3, 5, 9, 5, 6, 7), ncol = 3)
> A
```

	[,1]	[,2]	[,3]
[1,]	2	3	5
[2,]	2	5	6
[3,]	3	9	7

We want to find the matrix $B = A^{-1}$. To start, we append to A the identity matrix and call the result AB :

```
> AB <- cbind(A, diag(c(1, 1, 1)))
> AB
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	2	3	5	1	0	0
[2,]	2	5	6	0	1	0
[3,]	3	9	7	0	0	1

On a matrix we allow ourselves to do the following three operations (sometimes called elementary operations) as often as we want:

1. Multiply a row by a (non-zero) constant.
2. Multiply a row by a (non-zero) constant and add the result to another row.
3. Interchange two rows.

The aim is to perform such operations on AB in a way such that one ends up with a 3×6 matrix which has the identity matrix in the three leftmost columns. The three rightmost columns will then contain $B = A^{-1}$.

Recall that writing e.g. $AB[1,]$ extracts the entire first row of AB .

- First, we make sure that $AB[1,1]=1$. Then we subtract a constant times the first row from the second to obtain that $AB[2,1]=0$, and similarly for the third row:

```
> AB[1, ] <- AB[1, ]/AB[1, 1]
> AB[2, ] <- AB[2, ] - 2 * AB[1, ]
> AB[3, ] <- AB[3, ] - 3 * AB[1, ]
> AB
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	1	1.5	2.5	0.5	0	0
[2,]	0	2.0	1.0	-1.0	1	0
[3,]	0	4.5	-0.5	-1.5	0	1

- Next we ensure that $AB[2,2]=1$. Afterwards we subtract a constant times the second row from the third to obtain that $AB[3,2]=0$:

```
> AB[2, ] <- AB[2, ]/AB[2, 2]
> AB[3, ] <- AB[3, ] - 4.5 * AB[2, ]
```

- Now we rescale the third row such that $AB[3,3]=1$:

```
> AB[3, ] <- AB[3, ]/AB[3, 3]
> AB
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	1	1.5	2.5	0.5000000	0.0000000	0.0000000
[2,]	0	1.0	0.5	-0.5000000	0.5000000	0.0000000
[3,]	0	0.0	1.0	-0.2727273	0.8181818	-0.3636364

Then AB has zeros below the main diagonal.

- We then work our way up to obtain that AB has zeros above the main diagonal:

```
> AB[2, ] <- AB[2, ] - 0.5 * AB[3, ]
> AB[1, ] <- AB[1, ] - 2.5 * AB[3, ]
> AB
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	1	1.5	0	1.1818182	-2.04545455	0.9090909
[2,]	0	1.0	0	-0.3636364	0.09090909	0.1818182
[3,]	0	0.0	1	-0.2727273	0.81818182	-0.3636364

```
> AB[1, ] <- AB[1, ] - 1.5 * AB[2, ]
> AB
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	1	0	0	1.7272727	-2.18181818	0.6363636
[2,]	0	1	0	-0.3636364	0.09090909	0.1818182
[3,]	0	0	1	-0.2727273	0.81818182	-0.3636364

Now we extract the three rightmost columns of AB into the matrix B. We claim that B is the inverse of A, and this can be verified by a simple matrix multiplication

```
> B <- AB[, 4:6]
> A %*% B
```

	[,1]	[,2]	[,3]
[1,]	1.000000e+00	3.330669e-16	1.110223e-16
[2,]	-4.440892e-16	1.000000e+00	2.220446e-16
[3,]	-2.220446e-16	9.992007e-16	1.000000e+00

So, apart from rounding errors, the product is the identity matrix, and hence $B = A^{-1}$. This example illustrates that numerical precision and rounding errors is an important issue when making computer programs.

4 Least squares

Consider the table of pairs (x_i, y_i) below.

x	1.00	2.00	3.00	4.00	5.00
y	3.70	4.20	4.90	5.70	6.00

A plot of y_i against x_i is shown in Figure 6.




Figure 6: Regression

The plot in Figure 6 suggests an approximately linear relationship between y and x , i.e.

$$y_i = \beta_0 + \beta_1 x_i \text{ for } i = 1, \dots, 5$$

Writing this in matrix form gives

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \dots \\ y_5 \end{bmatrix} \approx \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_5 \end{bmatrix} \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix} = \mathbf{X}\boldsymbol{\beta}$$

The first question is: Can we find a vector $\boldsymbol{\beta}$ such that $y = \mathbf{X}\boldsymbol{\beta}$? The answer is clearly no, because that would require the points to lie exactly on a straight line.

A more modest question is: Can we find a vector $\hat{\boldsymbol{\beta}}$ such that $\mathbf{X}\hat{\boldsymbol{\beta}}$ is in a sense “as close to y as possible”. The answer is yes. The task is to find $\hat{\boldsymbol{\beta}}$ such that the length of the vector

$$e = y - \mathbf{X}\boldsymbol{\beta}$$

is as small as possible. The solution is

$$\hat{\boldsymbol{\beta}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top y$$

```
> y
```

```
[1] 3.7 4.2 4.9 5.7 6.0
```

```
> X
```

```
      x
[1,] 1 1
[2,] 1 2
[3,] 1 3
[4,] 1 4
[5,] 1 5
```



```
> beta.hat <- solve(t(X) %*% X) %*% t(X) %*% y
> beta.hat
```

```
 [,1]
 3.07
x 0.61
```

5 A neat little exercise – from a bird’s perspective

On a sunny day, two tables are standing in an English country garden. On each table birds of unknown species are sitting having the time of their lives.

A bird from the first table says to those on the second table: “Hi – if one of you come to our table then there will be the same number of us on each table”. “Yeah, right”, says a bird from the second table, “but if one of you comes to our table, then we will be twice as many on our table as on yours”.

Question: How many birds are on each table? More specifically,

- Write up two equations with two unknowns.
- Solve these equations using the methods you have learned from linear algebra.
- Simply finding the solution by trial-and-error is considered cheating.